

Project1

November 13, 2021

1 Mercedes Benz Greener Manufacturing

In this project our main target is to reduce the time a Mercedes-Benz spends on the test bench. First we import all required libraries and datasets. Then we read, understand, clean the data and after that we analyze it.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df_train=pd.read_csv("train.csv")
df_test=pd.read_csv("test.csv")
df_train.head()
```

```
[2]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8 ... X375 X376 X377 X378 X379 \
0   0  130.81   k  v   at  a  d  u  j  o ...    0    0    1    0    0
1   6   88.53   k  t   av  e  d  y  l  o ...    1    0    0    0    0
2   7   76.26  az  w   n  c  d  x  j  x ...    0    0    0    0    0
3   9   80.62  az  t   n  f  d  x  l  e ...    0    0    0    0    0
4  13   78.02  az  v   n  f  d  h  d  n ...    0    0    0    0    0
```

```
      X380 X382 X383 X384 X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
```

[5 rows x 378 columns]

Here 'y' is our target variable and the rest are independent variables.

```
[3]: df_test.head()
```

```
[3]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	\
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	0	0	0	0
3	0	0	0	0
4	0	0	0	0

[5 rows x 377 columns]

In both datasets 'ID' column is not required, so we drop it.

```
[4]: df_train.drop(['ID'],axis=1,inplace=True)
df_test.drop(['ID'],axis=1,inplace=True)
```

```
[5]: df_train.shape # 4209 rows and 377 columns in training data
```

```
[5]: (4209, 377)
```

```
[6]: df_test.shape
```

```
[6]: (4209, 376)
```

```
[7]: # Check if there are any null values
print(df_train.isnull().any().sum())
print(df_test.isnull().any().sum())
```

```
0
```

```
0
```

Both values are zero means there is no null value in both datasets.

```
[8]: # Look at basic statistics of train data.
df_train.describe().T
```

```
[8]:
```

	count	mean	std	min	25%	50%	75%	max
y	4209.0	100.669318	12.679381	72.11	90.82	99.15	109.01	265.32
X10	4209.0	0.013305	0.114590	0.00	0.00	0.00	0.00	1.00
X11	4209.0	0.000000	0.000000	0.00	0.00	0.00	0.00	0.00
X12	4209.0	0.075077	0.263547	0.00	0.00	0.00	0.00	1.00
X13	4209.0	0.057971	0.233716	0.00	0.00	0.00	0.00	1.00
...

X380	4209.0	0.008078	0.089524	0.00	0.00	0.00	0.00	1.00
X382	4209.0	0.007603	0.086872	0.00	0.00	0.00	0.00	1.00
X383	4209.0	0.001663	0.040752	0.00	0.00	0.00	0.00	1.00
X384	4209.0	0.000475	0.021796	0.00	0.00	0.00	0.00	1.00
X385	4209.0	0.001426	0.037734	0.00	0.00	0.00	0.00	1.00

[369 rows x 8 columns]

```
[9]: # Check correlation of training data
df_train.corr()
```

```
[9]:
```


	y	X10	X11	X12	X13	X14	X15	\
y	1.000000	-0.026985	NaN	0.089792	0.048276	0.193643	0.023116	
X10	-0.026985	1.000000	NaN	-0.033084	-0.028806	-0.100474	-0.002532	
X11	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
X12	0.089792	-0.033084	NaN	1.000000	0.214825	-0.246513	-0.006212	
X13	0.048276	-0.028806	NaN	0.214825	1.000000	-0.083141	-0.005409	
...	
X380	0.040932	-0.010479	NaN	-0.005566	0.023045	0.007743	-0.001968	
X382	-0.159815	-0.010164	NaN	-0.024937	-0.021713	0.012713	-0.001908	
X383	0.040291	-0.004740	NaN	-0.011628	-0.010125	0.023604	-0.000890	
X384	-0.004591	-0.002532	NaN	-0.006212	0.041242	0.025199	-0.000475	
X385	-0.022280	-0.004387	NaN	-0.010765	-0.009373	0.043667	-0.000824	
	X16	X17	X18	...	X375	X376	X377	\
y	0.048946	-0.159815	-0.001789	...	0.029100	0.114005	0.061403	
X10	-0.005944	-0.010164	-0.010323	...	0.165277	-0.028618	-0.074244	
X11	NaN	NaN	NaN	...	NaN	NaN	NaN	
X12	-0.014584	-0.024937	-0.025327	...	-0.107864	-0.070214	0.030134	
X13	-0.012698	-0.021713	-0.010525	...	-0.169721	-0.061136	0.357229	
...	
X380	-0.004619	-0.007899	-0.008022	...	-0.061741	-0.022240	-0.061168	
X382	-0.004480	1.000000	0.085256	...	-0.059883	-0.021571	-0.059327	
X383	-0.002089	-0.003572	0.062481	...	-0.015413	-0.010059	0.035107	
X384	-0.001116	-0.001908	-0.001938	...	-0.014917	-0.005373	0.008694	
X385	-0.001934	-0.003307	-0.003359	...	0.055225	-0.009311	-0.025610	
	X378	X379	X380	X382	X383	X384	X385	
y	-0.258679	0.067919	0.040932	-0.159815	0.040291	-0.004591	-0.022280	
X10	-0.016870	-0.011374	-0.010479	-0.010164	-0.004740	-0.002532	-0.004387	
X11	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
X12	-0.016043	-0.027907	-0.005566	-0.024937	-0.011628	-0.006212	-0.010765	
X13	-0.036040	-0.024299	0.023045	-0.021713	-0.010125	0.041242	-0.009373	
...	
X380	-0.013110	-0.008839	1.000000	-0.007899	-0.003683	-0.001968	-0.003410	
X382	-0.012716	-0.008573	-0.007899	1.000000	-0.003572	-0.001908	-0.003307	
X383	-0.005930	-0.003998	-0.003683	-0.003572	1.000000	-0.000890	-0.001542	

```
X384 -0.003168 -0.002136 -0.001968 -0.001908 -0.000890 1.000000 -0.000824
X385 -0.005489 -0.003701 -0.003410 -0.003307 -0.001542 -0.000824 1.000000
```

```
[369 rows x 369 columns]
```

```
[10]: # Now we split response and independent variables. There is no y_test value,
      ↪ as we have to predict it.
X_train=df_train.drop(['y'],axis=1)
y_train=df_train['y']
X_test=df_test.iloc[:,:]
```

```
[11]: # Distribution plot of y
sns.distplot(y_train)
plt.show()
```



In correlation matrix there are some nan values and also in description some columns have same max. and min. values that means some columns have constant values or only zero values, so we drop them.

```
[12]: # X_train data
a=list()
for i in X_train.columns:
    if X_train[i].min() == X_train[i].max():
        a.append(i)
print(a)
```

```
X_train.drop(a,axis=1,inplace=True)
```

```
['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297',  
'X330', 'X347']
```

```
[13]: # X_test data  
b=list()  
for i in X_test.columns:  
    if X_test[i].min() == X_test[i].max():  
        b.append(i)  
print(b)  
X_test.drop(a,axis=1,inplace=True)
```

```
['X257', 'X258', 'X295', 'X296', 'X369']
```

Now we encode object type columns 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8' as we do not want object type data.

```
[14]: # Check unique categories of object type data.  
print(X_train['X0'].unique())  
print(X_test['X0'].unique())
```

```
['k' 'az' 't' 'al' 'o' 'w' 'j' 'h' 's' 'n' 'ay' 'f' 'x' 'y' 'aj' 'ak' 'am'  
'z' 'q' 'at' 'ap' 'v' 'af' 'a' 'e' 'ai' 'd' 'aq' 'c' 'aa' 'ba' 'as' 'i'  
'r' 'b' 'ax' 'bc' 'u' 'ad' 'au' 'm' 'l' 'aw' 'ao' 'ac' 'g' 'ab']  
['az' 't' 'w' 'y' 'x' 'f' 'ap' 'o' 'ay' 'al' 'h' 'z' 'aj' 'd' 'v' 'ak'  
'ba' 'n' 'j' 's' 'af' 'ax' 'at' 'aq' 'av' 'm' 'k' 'a' 'e' 'ai' 'i' 'ag'  
'b' 'am' 'aw' 'as' 'r' 'ao' 'u' 'l' 'c' 'ad' 'au' 'bc' 'g' 'an' 'ae' 'p'  
'bb']
```

```
[15]: # Check their length is same or not.  
print(len(X_train['X0'].unique()))  
print(len(X_test['X0'].unique()))
```

```
47
```

```
49
```

```
[16]: # Check it for another column.  
print(len(X_train['X2'].unique()))  
print(len(X_test['X2'].unique()))
```

```
44
```

```
45
```

As object type columns have different categories so we have to encode them manually.

```
[17]: variable=['X0','X1','X2','X3','X4','X5','X6','X8']  
for i in variable:  
    d=list(X_train[i].unique())
```

```

f=list(X_test[i].unique())
for j in f:
    if j not in d:
        d.append(j)
enco=dict(zip(d,range(len(d))))
X_train[i]=X_train[i].replace(enco)
X_test[i]=X_test[i].replace(enco)

```

```
[18]: X_train.head()
```

```

[18]:   X0  X1  X2  X3  X4  X5  X6  X8  X10  X12  ...  X375  X376  X377  X378  \
0    0   0   0   0   0   0   0   0   0   0  ...    0     0     1     0
1    0   1   1   1   0   1   1   0   0   0  ...    1     0     0     0
2    1   2   2   2   0   2   0   1   0   0  ...    0     0     0     0
3    1   1   2   3   0   2   1   2   0   0  ...    0     0     0     0
4    1   0   2   3   0   3   2   3   0   0  ...    0     0     0     0

      X379  X380  X382  X383  X384  X385
0         0     0     0     0     0     0
1         0     0     0     0     0     0
2         0     0     1     0     0     0
3         0     0     0     0     0     0
4         0     0     0     0     0     0

```

[5 rows x 364 columns]

```
[19]: X_test.head()
```

```

[19]:   X0  X1  X2  X3  X4  X5  X6  X8  X10  X12  ...  X375  X376  X377  X378  \
0    1   0   2   3   0  29   5  16   0   0  ...    0     0     0     1
1    2   3   7   0   0  30   6  18   0   0  ...    0     0     1     0
2    1   0   4   3   0  31   0  13   0   0  ...    0     0     0     1
3    1   5   2   3   0  32   1   3   0   0  ...    0     0     0     1
4    5   6   4   2   0   1   4   8   0   0  ...    1     0     0     0

      X379  X380  X382  X383  X384  X385
0         0     0     0     0     0     0
1         0     0     0     0     0     0
2         0     0     0     0     0     0
3         0     0     0     0     0     0
4         0     0     0     0     0     0

```

[5 rows x 364 columns]

In both data there are too many columns so we apply Principal Component Analysis to extract important features.

```
[20]: from sklearn.decomposition import PCA
```


```
[21]: # First we plot PCA with all components and see where we get smooth curve
pca = PCA(n_components=X_train.shape[1])
pca.fit(X_train)
```

```
[21]: PCA(n_components=364)
```

```
[22]: pca.n_components_
```

```
[22]: 364
```

```
[23]: plt.figure(figsize = (15,6))
sns.lineplot(data=pca.explained_variance_)
plt.show()
# In this graph we get maximum information less than 25 components.
```




```
[24]: # Now check where we get 95% variance explained.
pca = PCA(n_components = 0.95)
X_train_pca = pca.fit_transform(X_train)
print(X_train_pca.shape)
```

```
(4209, 12)
```

```
[25]: X_test_pca = pca.transform(X_test)
print(X_test_pca.shape)
```

```
(4209, 12)
```

```
[26]: plt.figure(figsize = (15,6))
sns.lineplot(data=pca.explained_variance_)
plt.show()
```



Above graph n=5 component has a elbow shape curve which explained maximum variance. So, we reduce PCA to 5 components.

```
[27]: pca = PCA(n_components = 5)
      X_train_pca = pca.fit_transform(X_train)
      print(X_train_pca.shape)
```

```
(4209, 5)
```

```
[28]: X_test_pca = pca.transform(X_test)
      print(X_test_pca.shape)
```

```
(4209, 5)
```

Now data is clean, dimension is reduced and ready to fit in model. First we use eXtreme Gradient Boosting Regressor of ensemble method. Then we try to fit in another models.

```
[29]: import xgboost
      from sklearn.metrics import r2_score, mean_squared_error
```

```
[30]: xgb_model = xgboost.XGBRegressor(n_estimators=1000)
      xgb_model.fit(X_train_pca, y_train, eval_metric='rmse')
```

```
[30]: XGBRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
                  colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                  importance_type='gain', interaction_constraints=None,
                  learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                  min_child_weight=1, missing=nan, monotone_constraints=None,
                  n_estimators=1000, n_jobs=0, num_parallel_tree=1, random_state=0,
                  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
                  tree_method=None, validate_parameters=False, verbosity=None)
```



```
[31]: xgb_pred = xgb_model.predict(X_train_pca)
```

```
[32]: r2_score(y_pred = xgb_pred, y_true=y_train) # r square value
```

```
[32]: 0.9713683061811162
```

```
[33]: mean_squared_error(y_pred = xgb_pred, y_true=y_train) # mse
```

```
[33]: 4.601929593445253
```

```
[34]: mean_squared_error(y_pred = xgb_pred, y_true=y_train,squared=False) # rmse
```

```
[34]: 2.145210850579787
```

Here XGB model fits very well and R square value is 0.97 (very high) that means it explained 97% variance of model. Mean square error and Root mean square error is 4.6 and 2.1 respectively, which means low error.

```
[35]: # Now predict y_test
xgb_pred = xgb_model.predict(X_test_pca)
xgb_pred
```

```
[35]: array([ 75.20987,  87.23393,  80.63685, ..., 103.73881, 114.57336,
          93.11466], dtype=float32)
```

```
[36]: y_test=pd.DataFrame(X_test_pca,xgb_pred)
```

```
[37]: # Final result for X_test
y_test.head()
```

```
[37]:
```

	0	1	2	3	4
75.209869	-10.157275	7.308688	11.069408	-5.872025	-7.740356
87.233932	-8.533170	9.980976	11.856146	-1.535415	-4.083620
80.636848	-9.981792	5.078707	13.852080	-4.190252	-6.625144
77.718788	-10.507914	-4.100156	17.032831	-6.478453	-1.472961
110.676888	-8.827351	-6.798141	-14.186807	-2.166010	-0.546464

Now we try to fit the data in Linear Regression model.

```
[38]: from sklearn.linear_model import LinearRegression
```

```
[39]: lr_model=LinearRegression()
lr_model.fit(X_train,y_train)
```

```
[39]: LinearRegression()
```

```
[40]: lr_pred=lr_model.predict(X_train)
```

```
[41]: r2_score(y_pred = lr_pred, y_true=y_train) # r square value
```

```
[41]: 0.5919691916499641
```

```
[42]: mean_squared_error(y_pred = lr_pred, y_true=y_train) # mse
```

```
[42]: 65.58218538733398
```

Here Linear Regression model does not fit well as R square value is very low and MSE is very high.
Lets try Ridge, Lasso and ElasticNet model.

```
[43]: from sklearn.linear_model import Ridge, Lasso, ElasticNet
```

```
[44]: # Ridge regression  
ridge_model = Ridge(alpha=0.1)  
ridge_model.fit(X_train,y_train)
```

```
[44]: Ridge(alpha=0.1)
```

```
[45]: ridge_pred = ridge_model.predict(X_train)
```

```
[46]: r2_score(y_pred = ridge_pred, y_true=y_train) # r square value
```

```
[46]: 0.5918099470610183
```

```
[47]: mean_squared_error(y_pred = ridge_pred, y_true=y_train) # mse
```

```
[47]: 65.6077805334368
```

```
[48]: # Lasso regression  
lasso_model = Lasso(alpha=0.1)  
lasso_model.fit(X_train,y_train)
```

```
[48]: Lasso(alpha=0.1)
```

```
[49]: lasso_pred = ridge_model.predict(X_train)
```

```
[50]: r2_score(y_pred = lasso_pred, y_true=y_train) # r square value
```

```
[50]: 0.5918099470610183
```

```
[51]: mean_squared_error(y_pred = lasso_pred, y_true=y_train) # mse
```

```
[51]: 65.6077805334368
```

```
[52]: # ElasticNet Regression  
enet_model = ElasticNet(alpha=0.1, l1_ratio=0.5)
```

```
enet_model.fit(X_train,y_train)
enet_pred = enet_model.predict(X_train)
```

```
[53]: r2_score(y_pred = enet_pred, y_true=y_train) # r square value
```

```
[53]: 0.5380441696500797
```

```
[54]: mean_squared_error(y_pred = enet_pred, y_true=y_train) # mse
```

```
[54]: 74.24947402691784
```

Here Ridge, Lasso and ElasticNet model does not fit well as R square value is very low and MSE is very high like LR model.

Till now XGBoost performs best. Let us try XGBoost model with Grid Search Cross Validation.

```
[55]: from sklearn.model_selection import GridSearchCV
```

```
[56]: param_grid = {'C': [0.1,1], 'gamma': [1,0.1]}
```

```
[57]: xgb_grid=GridSearchCV(xgb_model,param_grid,refit=True,verbose=2)
```

```
[58]: xgb_grid.fit(X_train_pca,y_train)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
[CV] END ...C=0.1, gamma=1; total time= 6.2s
[CV] END ...C=0.1, gamma=1; total time= 6.0s
[CV] END ...C=0.1, gamma=1; total time= 6.9s
[CV] END ...C=0.1, gamma=1; total time= 6.3s
[CV] END ...C=0.1, gamma=1; total time= 5.9s
[CV] END ...C=0.1, gamma=0.1; total time= 6.9s
[CV] END ...C=0.1, gamma=0.1; total time= 7.4s
[CV] END ...C=0.1, gamma=0.1; total time= 6.2s
[CV] END ...C=0.1, gamma=0.1; total time= 6.7s
[CV] END ...C=0.1, gamma=0.1; total time= 6.3s
[CV] END ...C=1, gamma=1; total time= 6.9s
[CV] END ...C=1, gamma=1; total time= 6.3s
[CV] END ...C=1, gamma=1; total time= 6.9s
[CV] END ...C=1, gamma=1; total time= 6.7s
[CV] END ...C=1, gamma=1; total time= 5.9s
[CV] END ...C=1, gamma=0.1; total time= 6.3s
[CV] END ...C=1, gamma=0.1; total time= 6.3s
[CV] END ...C=1, gamma=0.1; total time= 6.2s
[CV] END ...C=1, gamma=0.1; total time= 6.1s
[CV] END ...C=1, gamma=0.1; total time= 6.6s
```

```
[58]: GridSearchCV(estimator=XGBRegressor(base_score=0.5, booster=None,
                                         colsample_bylevel=1, colsample_bynode=1,
```

```

        colsample_bytree=1, gamma=0, gpu_id=-1,
        importance_type='gain',
        interaction_constraints=None,
        learning_rate=0.300000012, max_delta_step=0,
        max_depth=6, min_child_weight=1,
        missing=nan, monotone_constraints=None,
        n_estimators=1000, n_jobs=0,
        num_parallel_tree=1, random_state=0,
        reg_alpha=0, reg_lambda=1,
        scale_pos_weight=1, subsample=1,
        tree_method=None, validate_parameters=False,
        verbosity=None),
    param_grid={'C': [0.1, 1], 'gamma': [1, 0.1]}, verbose=2)

```

```
[59]: grid_predictions = xgb_grid.predict(X_train_pca)
```

```
[60]: r2_score(y_train,grid_predictions)
```

```
[60]: 0.9664775171018932
```

```
[61]: mean_squared_error(y_pred = grid_predictions, y_true=y_train) # mse
```

```
[61]: 5.3880188531778055
```

```
[62]: mean_squared_error(y_pred = grid_predictions, y_true=y_train,squared=False) #RMSE
      ↪rmse
```

```
[62]: 2.3212106438619062
```

This model also performs well as R square value is 0.96 (very high) that means it explained 96.6% variance of model. Mean square error and Root mean square error is 5.3 and 2.3 respectively, which means low error.

```
[63]: # Predict y_test with this model.
      grid_pred = xgb_grid.predict(X_test_pca)
      grid_pred
```

```
[63]: array([ 80.864975,  90.027245,  82.5852   , ..., 101.24793 , 115.40703 ,
          95.76787 ], dtype=float32)
```

```
[65]: y_test2=pd.DataFrame(X_test_pca,grid_pred)
```

```
[66]: y_test2.head() # final result with Grid search cv
```

```
[66]:
```

	0	1	2	3	4
80.864975	-10.157275	7.308688	11.069408	-5.872025	-7.740356
90.027245	-8.533170	9.980976	11.856146	-1.535415	-4.083620
82.585197	-9.981792	5.078707	13.852080	-4.190252	-6.625144

```
75.685059 -10.507914 -4.100156 17.032831 -6.478453 -1.472961
107.953415 -8.827351 -6.798141 -14.186807 -2.166010 -0.546464
```

```
[67]: y_test.head()           # final result without Grid search cv
```

```
[67]:
```

	0	1	2	3	4
75.209869	-10.157275	7.308688	11.069408	-5.872025	-7.740356
87.233932	-8.533170	9.980976	11.856146	-1.535415	-4.083620
80.636848	-9.981792	5.078707	13.852080	-4.190252	-6.625144
77.718788	-10.507914	-4.100156	17.032831	-6.478453	-1.472961
110.676888	-8.827351	-6.798141	-14.186807	-2.166010	-0.546464

Both results are very good we can use any one of them but without Grid search result is better from with grid search result.

```
[ ]: # END
```